

FC Net

In this section all relevant python code files will be displayed in their entirety. Each file will be labeled accordingly.

fc_net.py

```
import numpy as np

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of  $D$ , a hidden dimension of  $H$ , and perform classification over  $C$  classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead, it
    will interact with a separate Solver object that is responsible for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
                  dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg
```

```

# ===== #
# YOUR CODE HERE:
#   Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
#   self.params['W2'], self.params['b1'] and self.params['b2']. The
#   biases are initialized to zero and the weights are initialized
#   so that each parameter has mean 0 and standard deviation weight_scale.
#   The dimensions of W1 should be (input_dim, hidden_dim) and the
#   dimensions of W2 should be (hidden_dims, num_classes)
# ===== #

self.params['W1'] = np.random.randn(input_dim,hidden_dims)*weight_scale
self.params['b1'] = np.zeros((hidden_dims,))
self.params['W2'] = np.random.randn(hidden_dims,num_classes)*weight_scale
self.params['b2'] = np.zeros((num_classes,))
# ===== #
# END YOUR CODE HERE
# ===== #

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and return:
    - scores: Array of shape (N, C) giving classification scores, where
      scores[i, c] is the classification score for X[i] and class c.

    If y is not None, then run a training-time forward and backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping parameter
      names to gradients of the loss with respect to those parameters.
    """
    scores = None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the forward pass of the two-layer neural network. Store
    #   the class scores as the variable 'scores'. Be sure to use the layers
    #   you prior implemented.
    # ===== #

```

```

h1,cache1 = affine_relu_forward(X,self.params['W1'],self.params['b1'])
scores,cache2 = affine_forward(h1,self.params['W2'],self.params['b2'])
# ===== #
# END YOUR CODE HERE
# ===== #

# If y is None then we are in test mode so just return scores
if y is None:
    return scores

loss, grads = 0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backward pass of the two-layer neural net. Store
# the loss as the variable 'loss' and store the gradients in the
# 'grads' dictionary. For the grads dictionary, grads['W1'] holds
# the gradient for W1, grads['b1'] holds the gradient for b1, etc.
# i.e., grads[k] holds the gradient for self.params[k].
#
# Add L2 regularization, where there is an added cost  $0.5 \cdot \text{self.reg} \cdot W^2$ 
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# ===== #

loss,dLdz = softmax_loss(scores,y)
loss += 0.5*self.reg*np.linalg.norm(self.params['W1'])**2 + 0.5*self.reg*np.linalg.norm(
dLdh1,dLdw2, dLdb2 = affine_backward(dLdz,cache2)
dLdx,dLdw1,dLdb1 = affine_relu_backward(dLdh1,cache1)
grads['W2'] = dLdw2 + self.reg*self.params['W2']
grads['b2'] = dLdb2
grads['W1'] = dLdw1 + self.reg*self.params['W1']
grads['b1'] = dLdb1
# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

```

class FullyConnectedNet(object):

```

```

    """

```

A fully-connected neural network with an arbitrary number of hidden layers, ReLU nonlinearities, and a softmax loss function. This will also implement dropout and batch normalization as options. For a network with L layers,

the architecture will be

$\{ \text{affine} - [\text{batch norm}] - \text{relu} - [\text{dropout}] \} \times (L - 1) - \text{affine} - \text{softmax}$

where batch normalization and dropout are optional, and the $\{...\}$ block is repeated $L - 1$ times.

Similar to the `TwoLayerNet` above, learnable parameters are stored in the `self.params` dictionary and will be learned using the `Solver` class.

"""

```
def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=0, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):
```

"""

Initialize a new `FullyConnectedNet`.

Inputs:

- `hidden_dims`: A list of integers giving the size of each hidden layer.
- `input_dim`: An integer giving the size of the input.
- `num_classes`: An integer giving the number of classes to classify.
- `dropout`: Scalar between 0 and 1 giving dropout strength. If `dropout=0` then the network should not use dropout at all.
- `use_batchnorm`: Whether or not the network should use batch normalization.
- `reg`: Scalar giving L2 regularization strength.
- `weight_scale`: Scalar giving the standard deviation for random initialization of the weights.
- `dtype`: A numpy datatype object; all computations will be performed using this datatype. `float32` is faster but less accurate, so you should use `float64` for numeric gradient checking.
- `seed`: If not `None`, then pass this random seed to the dropout layers. This will make the dropout layers deterministic so we can gradient check the model.

"""

```
self.use_batchnorm = use_batchnorm
self.use_dropout = dropout > 0
self.reg = reg
self.num_layers = 1 + len(hidden_dims)
self.dtype = dtype
self.params = {}
```

```
# ===== #
```

```
# YOUR CODE HERE:
```

```
# Initialize all parameters of the network in the self.params dictionary.
# The weights and biases of layer 1 are  $W_1$  and  $b_1$ ; and in general the
# weights and biases of layer  $i$  are  $W_i$  and  $b_i$ . The
```

```

#   biases are initialized to zero and the weights are initialized
#   so that each parameter has mean 0 and standard deviation weight_scale.
#   ===== #

for i in range(self.num_layers):
    w_i = 'W'+str(i+1)
    b_i = 'b'+str(i+1)
    if(i==0):
        self.params[w_i] = np.random.randn(input_dim,hidden_dims[i])*weight_scale
        self.params[b_i] = np.zeros((hidden_dims[i],))
    elif(i==self.num_layers-1):
        self.params[w_i] = np.random.randn(hidden_dims[i-1],num_classes)*weight_scale
        self.params[b_i] = np.zeros((num_classes,))
    else:
        self.params[w_i] = np.random.randn(hidden_dims[i-1],hidden_dims[i])*weight_scale
        self.params[b_i] = np.zeros((hidden_dims[i],))

#   ===== #
#   END YOUR CODE HERE
#   ===== #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
    if seed is not None:
        self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """

```

Compute loss and gradient for the fully-connected net.

Input / output: Same as TwoLayerNet above.

"""

X = X.astype(self.dtype)

mode = 'test' if y is None else 'train'

*# Set train/test mode for batchnorm params and dropout param since they
behave differently during training and testing.*

if self.dropout_param is not None:
 self.dropout_param['mode'] = mode

if self.use_batchnorm:
 for bn_param in self.bn_params:
 bn_param[mode] = mode

scores = None

=====

YOUR CODE HERE:

*# Implement the forward pass of the FC net and store the output
scores as the variable "scores".*

=====

hs = []

caches= []

for i in range(self.num_layers):

#generate keys

w_i = 'W'+str(i+1)

b_i = 'b'+str(i+1)

if(i==0):

h_i, cache = affine_relu_forward(X,self.params[w_i], self.params[b_i])

hs.append(h_i)

caches.append(cache)

elif(i!= self.num_layers-1):

h_i, cache = affine_relu_forward(hs[i-1],self.params[w_i], self.params[b_i])

hs.append(h_i)

caches.append(cache)

else:

h_i,cache = affine_forward(hs[i-1],self.params[w_i], self.params[b_i])

scores = h_i

caches.append(cache)

=====

END YOUR CODE HERE

=====

If test mode return early

if mode == 'test':

```

    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
#   Implement the backwards pass of the FC net and store the gradients
#   in the grads dict, so that grads[k] is the gradient of self.params[k]
#   Be sure your L2 regularization includes a 0.5 factor.
# ===== #
loss, dLdupstream = softmax_loss(scores, y)
dLdhi, dLdwi, dLdbi = 0, 0, 0
for i in reversed(range(self.num_layers)):
    #generate keys
    w_i = 'W'+str(i+1)
    b_i = 'b'+str(i+1)
    loss += 0.5*self.reg*np.linalg.norm(self.params[w_i])**2
    if(i == self.num_layers-1):
        dLdhi, dLdwi, dLdbi = affine_backward(dLdupstream, caches[i])
    else:
        dLdhi, dLdwi, dLdbi = affine_relu_backward(dLdupstream, caches[i])
    dLdupstream = dLdhi
    grads[w_i] = dLdwi + self.reg*self.params[w_i]
    grads[b_i] = dLdbi

# ===== #
# END YOUR CODE HERE
# ===== #
return loss, grads

```

layers.py

```

import numpy as np
import pdb

```

```

def affine_forward(x, w, b):
    """
    Computes the forward pass for an affine (fully-connected) layer.

    The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
    examples, where each example x[i] has shape (d_1, ..., d_k). We will
    reshape each input into a vector of dimension D = d_1 * ... * d_k, and
    then transform it to an output vector of dimension M.
    """

```

```

Inputs:
- x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
- w: A numpy array of weights, of shape (D, M)
- b: A numpy array of biases, of shape (M,)

Returns a tuple of:
- out: output, of shape (N, M)
- cache: (x, w, b)
"""

# ===== #
# YOUR CODE HERE:
#   Calculate the output of the forward pass. Notice the dimensions
#   of w are D x M, which is the transpose of what we did in earlier
#   assignments.
# ===== #
x_shapes = x.shape
D = np.prod(x_shapes[1:])
N = x_shapes[0]
X = x.reshape(N,D)
out = X@w+b

# ===== #
# END YOUR CODE HERE
# ===== #

cache = (x, w, b)
return out, cache

def affine_backward(dout, cache):
    """
    Computes the backward pass for an affine layer.

    Inputs:
    - dout: Upstream derivative, of shape (N, M)
    - cache: Tuple of:
      - x: Input data, of shape (N, d_1, ... d_k)
      - w: Weights, of shape (D, M)

    Returns a tuple of:
    - dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
    - dw: Gradient with respect to w, of shape (D, M)
    - db: Gradient with respect to b, of shape (M,)
    """

```



```

x, w, b = cache
dx, dw, db = None, None, None

# ===== #
# YOUR CODE HERE:
#   Calculate the gradients for the backward pass.
# ===== #

# dout is N x M
# dx should be N x d1 x ... x dk; it relates to dout through multiplication with w, which
# dw should be D x M; it relates to dout through multiplication with x, which is N x D af
# db should be M; it is just the sum over dout examples
x_shapes = x.shape
D = np.prod(x_shapes[1:])
N = x_shapes[0]
x_reshape = x.reshape(N,D)
dx = (dout@w.T).reshape(x_shapes)
dw = x_reshape.T@dout
db = dout.T@np.ones((N,))
# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dw, db

def relu_forward(x):
    """
    Computes the forward pass for a layer of rectified linear units (ReLU).

    Input:
    - x: Inputs, of any shape

    Returns a tuple of:
    - out: Output, of the same shape as x
    - cache: x
    """
    # ===== #
    # YOUR CODE HERE:
    #   Implement the ReLU forward pass.
    # ===== #

    out = (x>0)*x
    # ===== #
    # END YOUR CODE HERE
    # ===== #

```

```

    cache = x
    return out, cache

def relu_backward(dout, cache):
    """
    Computes the backward pass for a layer of rectified linear units (ReLU).

    Input:
    - dout: Upstream derivatives, of any shape
    - cache: Input x, of same shape as dout

    Returns:
    - dx: Gradient with respect to x
    """
    x = cache

    # ===== #
    # YOUR CODE HERE:
    #   Implement the ReLU backward pass
    # ===== #

    # ReLU directs linearly to those > 0
    dx = dout*(x>0)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return dx

def svm_loss(x, y):
    """
    Computes the loss and gradient using for multiclass SVM classification.

    Inputs:
    - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
        for the ith input.
    - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
        0 <= y[i] < C

    Returns a tuple of:
    - loss: Scalar giving the loss
    - dx: Gradient of the loss with respect to x
    """
    N = x.shape[0]

```

```

correct_class_scores = x[np.arange(N), y]
margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
margins[np.arange(N), y] = 0
loss = np.sum(margins) / N
num_pos = np.sum(margins > 0, axis=1)
dx = np.zeros_like(x)
dx[margins > 0] = 1
dx[np.arange(N), y] -= num_pos
dx /= N
return loss, dx

```



```

def softmax_loss(x, y):
    """
    Computes the loss and gradient for softmax classification.

    Inputs:
    - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
      for the ith input.
    - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
      0 <= y[i] < C

    Returns a tuple of:
    - loss: Scalar giving the loss
    - dx: Gradient of the loss with respect to x
    """
    probs = np.exp(x - np.max(x, axis=1, keepdims=True))
    probs /= np.sum(probs, axis=1, keepdims=True)
    N = x.shape[0]
    loss = -np.sum(np.log(probs[np.arange(N), y])) / N
    dx = probs.copy()
    dx[np.arange(N), y] -= 1
    dx /= N
    return loss, dx

```