

Two Layer NN

In this section all relevant python code files will be displayed in their entirety. Each file will be labeled accordingly.

neural_net.py

```
from optparse import IndentedHelpFormatter
from nndl.optim import sgd
import numpy as np
import matplotlib.pyplot as plt

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network. The net has an input dimension of
     $N$ , a hidden layer dimension of  $H$ , and performs classification over  $C$  classes.
    We train the network with a softmax loss function and L2 regularization on the
    weight matrices. The network uses a ReLU nonlinearity after the first fully
    connected layer.

    In other words, the network has the following architecture:

    input - fully connected layer - ReLU - fully connected layer - softmax

    The outputs of the second fully-connected layer are the scores for each class.
    """

    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
        """
        Initialize the model. Weights are initialized to small random values and
        biases are initialized to zero. Weights and biases are stored in the
        variable self.params, which is a dictionary with the following keys:

        W1: First layer weights; has shape (H, D)
        b1: First layer biases; has shape (H,)
        W2: Second layer weights; has shape (C, H)
        b2: Second layer biases; has shape (C,)

        Inputs:
        - input_size: The dimension  $D$  of the input data.
        - hidden_size: The number of neurons  $H$  in the hidden layer.
        - output_size: The number of classes  $C$ .
        """
        self.params = {}
```

```

self.params['W1'] = std * np.random.randn(hidden_size, input_size)
self.params['b1'] = np.zeros(hidden_size)
self.params['W2'] = std * np.random.randn(output_size, hidden_size)
self.params['b2'] = np.zeros(output_size)

def loss(self, X, y=None, reg=0.0):
    """
    Compute the loss and gradients for a two layer fully connected neural
    network.

    Inputs:
    - X: Input data of shape (N, D). Each X[i] is a training sample.
    - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
        an integer in the range 0 <= y[i] < C. This parameter is optional; if it
        is not passed then we only return scores, and if it is passed then we
        instead return the loss and gradients.
    - reg: Regularization strength.

    Returns:
    If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
    the score for class c on input X[i].

    If y is not None, instead return a tuple of:
    - loss: Loss (data loss and regularization loss) for this batch of training
        samples.
    - grads: Dictionary mapping parameter names to gradients of those parameters
        with respect to the loss function; has the same keys as self.params.
    """
    # Unpack variables from the params dictionary
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    N, D = X.shape

    # Compute the forward pass
    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Calculate the output scores of the neural network. The result
    # should be (N, C). As stated in the description for this class,
    # there should not be a ReLU layer after the second FC layer.
    # The output of the second FC layer is the output scores. Do not
    # use a for loop in your implementation.
    # ===== #
    r = W1@X.T+b1.reshape(len(b1),1)

```

```

h1 = (r>0)*r
scores = W2@h1+b2.reshape(len(b2),1)
scores = scores.T
# ===== #
# END YOUR CODE HERE
# ===== #

# If the targets are not given then jump out, we're done
if y is None:
    return scores

# Compute the loss
loss = None

# ===== #
# YOUR CODE HERE:
# Calculate the loss of the neural network. This includes the
# softmax loss and the L2 regularization for W1 and W2. Store the
# total loss in teh variable loss. Multiply the regularization
# loss by 0.5 (in addition to the factor reg).
# ===== #

# scores is num_examples by num_classes
C = scores.shape[1]
N = scores.shape[0]
escores = np.exp(scores)
loss = np.ones((1,N))@(np.log(escores@np.ones((C,))) - scores[np.arange(N),y])/N + reg*(
# ===== #
# END YOUR CODE HERE
# ===== #

grads = {}

# ===== #
# YOUR CODE HERE:
# Implement the backward pass. Compute the derivatives of the
# weights and the biases. Store the results in the grads
# dictionary. e.g., grads['W1'] should store the gradient for
# W1, and be of the same size as W1.
# ===== #
scaling = escores.T/(escores@np.ones((C,)))
scaling[y,np.arange(N)] -= 1
dLdCE = scaling@np.ones((N,))/N
#print(dLdCE.shape)
grads['b2'] = dLdCE

```

```

grads['W2'] = scaling@h1.T/N +reg*W2
dLdh1 = W2.T@dLdCE
dLdr = W2.T@scaling*(r>0)@np.ones((N,))/N
grads['b1'] = dLdr
grads['W1'] = W2.T@scaling*(r>0)@X/N +reg*W1
# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

def train(self, X, y, X_val, y_val,
          learning_rate=1e-3, learning_rate_decay=0.95,
          reg=1e-5, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
        X[i] has label c, where 0 <= c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the learning rate
        after each epoch.
    - reg: Scalar giving regularization strength.
    - num_iters: Number of steps to take when optimizing.
    - batch_size: Number of training examples to use per step.
    - verbose: boolean; if true print progress during optimization.
    """
    num_train = X.shape[0]
    iterations_per_epoch = max(num_train / batch_size, 1)

    # Use SGD to optimize the parameters in self.model
    loss_history = []
    train_acc_history = []
    val_acc_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:

```

```

# Create a minibatch by sampling batch_size samples randomly.
# ===== #
indexes = (np.random.rand(batch_size,1)*num_train).astype(int)
X_batch = X[indexes,:].reshape(batch_size,X.shape[1])
y_batch = y[indexes].reshape(batch_size,)
# ===== #
# END YOUR CODE HERE
# ===== #

# Compute loss and gradients using the current minibatch
loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
# Perform a gradient descent step using the minibatch to update
# all parameters (i.e., W1, W2, b1, and b2).
# ===== #
config={}
config['learning_rate'] = learning_rate
self.params['W1'],config = sgd(self.params['W1'],grads['W1'],config)
self.params['b1'],config = sgd(self.params['b1'],grads['b1'],config)
self.params['W2'],config = sgd(self.params['W2'],grads['W2'],config)
self.params['b2'],config = sgd(self.params['b2'],grads['b2'],config)
# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

# Every epoch, check train and val accuracy and decay learning rate.
if it % iterations_per_epoch == 0:
    # Check accuracy
    train_acc = (self.predict(X_batch) == y_batch).mean()
    val_acc = (self.predict(X_val) == y_val).mean()
    train_acc_history.append(train_acc)
    val_acc_history.append(val_acc)

    # Decay learning rate
    learning_rate *= learning_rate_decay

return {
    'loss_history': loss_history,
    'train_acc_history': train_acc_history,
    'val_acc_history': val_acc_history,

```

```

}

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict labels for
    data points. For each data point we predict scores for each of the C
    classes, and assign each data point to the class with the highest score.

    Inputs:
    - X: A numpy array of shape (N, D) giving N D-dimensional data points to
        classify.

    Returns:
    - y_pred: A numpy array of shape (N,) giving predicted labels for each of
        the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
        to have class c, where 0 <= c < C.
    """
    y_pred = None

    # ===== #
    # YOUR CODE HERE:
    #   Predict the class given the input data.
    # ===== #
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    N, D = X.shape
    r = W1@X.T+b1.reshape(len(b1),1)
    h1 = (r>0)*r
    scores = W2@h1+b2.reshape(len(b2),1)
    scores = scores.T
    y_pred = np.argmax(scores, axis=1)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```