

## fc\_net.py

This file contains all the code for fc\_net.py

## Code

```
import numpy as np
from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of  $D$ , a hidden dimension of  $H$ , and perform classification over  $C$  classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead, it
    will interact with a separate Solver object that is responsible for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
                  dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg

        # ===== #
```

```

# YOUR CODE HERE:
# Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
# self.params['W2'], self.params['b1'] and self.params['b2']. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation weight_scale.
# The dimensions of W1 should be (input_dim, hidden_dim) and the
# dimensions of W2 should be (hidden_dims, num_classes)
# ===== #

# ===== #
# END YOUR CODE HERE
# ===== #

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and return:
    - scores: Array of shape (N, C) giving classification scores, where
      scores[i, c] is the classification score for X[i] and class c.

    If y is not None, then run a training-time forward and backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping parameter
      names to gradients of the loss with respect to those parameters.
    """
    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the two-layer neural network. Store
    # the class scores as the variable 'scores'. Be sure to use the layers
    # you prior implemented.
    # ===== #

    # ===== #
    # END YOUR CODE HERE
    # ===== #

```

```

# If y is None then we are in test mode so just return scores
if y is None:
    return scores

loss, grads = 0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backward pass of the two-layer neural net. Store
# the loss as the variable 'loss' and store the gradients in the
# 'grads' dictionary. For the grads dictionary, grads['W1'] holds
# the gradient for W1, grads['b1'] holds the gradient for b1, etc.
# i.e., grads[k] holds the gradient for self.params[k].
#
# Add L2 regularization, where there is an added cost  $0.5 * \text{self.reg} * W^2$ 
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# ===== #

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

```

class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

    where batch normalization and dropout are optional, and the {...} block is
    repeated L - 1 times.

    Similar to the TwoLayerNet above, learnable parameters are stored in the
    self.params dictionary and will be learned using the Solver class.
    """

    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
                 dropout=0, use_batchnorm=False, reg=0.0,
                 weight_scale=1e-2, dtype=np.float32, seed=None):

```

```

"""
Initialize a new FullyConnectedNet.

Inputs:
- hidden_dims: A list of integers giving the size of each hidden layer.
- input_dim: An integer giving the size of the input.
- num_classes: An integer giving the number of classes to classify.
- dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
  the network should not use dropout at all.
- use_batchnorm: Whether or not the network should use batch normalization.
- reg: Scalar giving L2 regularization strength.
- weight_scale: Scalar giving the standard deviation for random
  initialization of the weights.
- dtype: A numpy datatype object; all computations will be performed using
  this datatype. float32 is faster but less accurate, so you should use
  float64 for numeric gradient checking.
- seed: If not None, then pass this random seed to the dropout layers. This
  will make the dropout layers deterministic so we can gradient check the
  model.
"""

self.use_batchnorm = use_batchnorm
self.use_dropout = dropout > 0
self.reg = reg
self.num_layers = 1 + len(hidden_dims)
self.dtype = dtype
self.params = {}

# ===== #
# YOUR CODE HERE:
#   Initialize all parameters of the network in the self.params dictionary.
#   The weights and biases of layer 1 are W1 and b1; and in general the
#   weights and biases of layer i are Wi and bi. The
#   biases are initialized to zero and the weights are initialized
#   so that each parameter has mean 0 and standard deviation weight_scale.
#
#   BATCHNORM: Initialize the gammas of each layer to 1 and the beta
#   parameters to zero. The gamma and beta parameters for layer 1 should
#   be self.params['gamma1'] and self.params['beta1']. For layer 2, they
#   should be gamma2 and beta2, etc. Only use batchnorm if self.use_batchnorm
#   is true and DO NOT do batch normalize the output scores.
# ===== #
for i in range(self.num_layers):
    w_i = 'W'+str(i+1)
    b_i = 'b'+str(i+1)
    gamma_i = 'gamma'+str(i+1)
    beta_i = 'beta'+str(i+1)

```

```

if(i==0):
    self.params[w_i] = np.random.randn(input_dim,hidden_dims[i])*weight_scale
    self.params[b_i] = np.zeros((hidden_dims[i],))
    if(self.use_batchnorm):
        self.params[gamma_i] = np.ones((hidden_dims[i],))
        self.params[beta_i] = np.zeros((hidden_dims[i],))
elif(i==self.num_layers-1):
    self.params[w_i] = np.random.randn(hidden_dims[i-1],num_classes)*weight_scale
    self.params[b_i] = np.zeros((num_classes,))
else:
    self.params[w_i] = np.random.randn(hidden_dims[i-1],hidden_dims[i])*weight_scale
    self.params[b_i] = np.zeros((hidden_dims[i],))
    if(self.use_batchnorm):
        self.params[gamma_i] = np.ones((hidden_dims[i],))
        self.params[beta_i] = np.zeros((hidden_dims[i],))

# ===== #
# END YOUR CODE HERE
# ===== #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
    if seed is not None:
        self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """

```

*Compute loss and gradient for the fully-connected net.*

*Input / output: Same as TwoLayerNet above.*

"""

X = X.astype(self.dtype)

mode = 'test' if y is None else 'train'

*# Set train/test mode for batchnorm params and dropout param since they  
# behave differently during training and testing.*

if self.dropout\_param is not None:

self.dropout\_param['mode'] = mode

if self.use\_batchnorm:

for bn\_param in self.bn\_params:

bn\_param['mode'] = mode

scores = None

# ===== #

# YOUR CODE HERE:

*# Implement the forward pass of the FC net and store the output  
# scores as the variable "scores".*

#

*# BATCHNORM: If self.use\_batchnorm is true, insert a bathnorm layer  
# between the affine\_forward and relu\_forward layers. You may  
# also write an affine\_batchnorm\_relu() function in layer\_utils.py.*

#

*# DROPOUT: If dropout is non-zero, insert a dropout layer after  
# every ReLU layer.*

# ===== #

hs = []

caches= []

dro\_caches = []

for i in range(self.num\_layers):

*#generate keys*

w\_i = 'W'+str(i+1)

b\_i = 'b'+str(i+1)

gamma\_i = 'gamma'+str(i+1)

beta\_i = 'beta'+str(i+1)

if(i==0):

if(self.use\_batchnorm):

h\_i, cache = affine\_batchnorm\_relu\_forward(X,self.params[w\_i], self.params[b\_i], s

if self.use\_dropout:

h\_i, dro\_cache = dropout\_forward(h\_i,self.dropout\_param)

dro\_caches.append(dro\_cache)

hs.append(h\_i)

caches.append(cache)

```

else:
    h_i, cache = affine_relu_forward(X, self.params[w_i], self.params[b_i])
    if self.use_dropout:
        h_i, dro_cache = dropout_forward(h_i, self.dropout_param)
        dro_caches.append(dro_cache)
    hs.append(h_i)
    caches.append(cache)

elif(i!= self.num_layers-1):
    if(self.use_batchnorm):
        h_i, cache = affine_batchnorm_relu_forward(hs[i-1], self.params[w_i], self.params[b_i])
        if self.use_dropout:
            h_i, dro_cache = dropout_forward(h_i, self.dropout_param)
            dro_caches.append(dro_cache)
        hs.append(h_i)
        caches.append(cache)

    else:
        h_i, cache = affine_relu_forward(hs[i-1], self.params[w_i], self.params[b_i])
        if self.use_dropout:
            h_i, dro_cache = dropout_forward(h_i, self.dropout_param)
            dro_caches.append(dro_cache)
        hs.append(h_i)
        caches.append(cache)

else:
    h_i, cache = affine_forward(hs[i-1], self.params[w_i], self.params[b_i])
    scores = h_i
    caches.append(cache)

# ===== #
# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
#
# BATCHNORM: Incorporate the backward pass of the batchnorm.
#

```

```

# DROPOUT: Incorporate the backward pass of dropout.
# ===== #
loss,dLdupstream = softmax_loss(scores,y)
dLdhi,dLdwi, dLdbi, dLdgammai, dLdbetai =0,0,0,0,0
for i in reversed(range(self.num_layers)):
    #generate keys
    w_i = 'W'+str(i+1)
    b_i = 'b'+str(i+1)
    gamma_i = 'gamma'+str(i+1)
    beta_i = 'beta'+str(i+1)
    loss += 0.5*self.reg*np.linalg.norm(self.params[w_i])**2
    if(i == self.num_layers-1):
        dLdhi,dLdwi, dLdbi = affine_backward(dLdupstream,caches[i])
    else:
        if(self.use_batchnorm):
            if self.use_dropout:
                dLdupstream = dropout_backward(dLdupstream,dro_caches[i])
                dLdhi,dLdwi,dLdbi, dLdgammai, dLdbetai = affine_batchnorm_relu_backward(dLdupstream,caches[i])
                grads[gamma_i] = dLdgammai
                grads[beta_i] = dLdbetai
            else:
                if self.use_dropout:
                    dLdupstream = dropout_backward(dLdupstream,dro_caches[i])
                    dLdhi,dLdwi,dLdbi = affine_relu_backward(dLdupstream,caches[i])
                dLdupstream = dLdhi
            grads[w_i]=dLdwi + self.reg*self.params[w_i]
            grads[b_i]=dLdbi

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```