

conv_layers.py

This file contains all the code in conv_layers.py

Code

```
from multiprocessing import pool
import numpy as np
from nndl.layers import *
import pdb


def conv_forward_naive(x, w, b, conv_param):
    """
    A naive implementation of the forward pass for a convolutional layer.

    The input consists of  $N$  data points, each with  $C$  channels, height  $H$  and width  $W$ . We convolve each input with  $F$  different filters, where each filter spans all  $C$  channels and has height  $HH$  and width  $WW$ .

    Input:
    -  $x$ : Input data of shape  $(N, C, H, W)$ 
    -  $w$ : Filter weights of shape  $(F, C, HH, WW)$ 
    -  $b$ : Biases, of shape  $(F,)$ 
    -  $conv\_param$ : A dictionary with the following keys:
      - 'stride': The number of pixels between adjacent receptive fields in the horizontal and vertical directions.
      - 'pad': The number of pixels that will be used to zero-pad the input.

    Returns a tuple of:
    -  $out$ : Output data, of shape  $(N, F, H', W')$  where  $H'$  and  $W'$  are given by
       $H' = 1 + (H + 2 * pad - HH) / stride$ 
       $W' = 1 + (W + 2 * pad - WW) / stride$ 
    -  $cache$ :  $(x, w, b, conv\_param)$ 
    """
    out = None
    pad = conv_param['pad']
    stride = conv_param['stride']

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of a convolutional neural network.
    # Store the output as 'out'.
    # Hint: to pad the array, you can use the function np.pad.
    # ===== #
```

```

npad = ((0,0), (0,0), (pad,pad), (pad,pad))
xpad = np.pad(x, npad, mode='constant', constant_values=0)
N, C, H, W = xpad.shape
F, C, HH, WW = w.shape
H_out = int(1 + (H - HH) / stride)
W_out = int(1 + (W - WW) / stride)
out = np.zeros((N, F, H_out, W_out))
for n in range(N):
    for i in range(F):
        for j in range(H_out):
            for k in range(W_out):
                out[n, i, j, k] = np.sum(w[i, :] * xpad[n, :, stride*j:stride*j+HH, stride*k:stride*k+WW])

# =====j== #
# END YOUR CODE HERE
# ===== #

cache = (x, w, b, conv_param)
return out, cache

def conv_backward_naive(dout, cache):
    """
    A naive implementation of the backward pass for a convolutional layer.

    Inputs:
    - dout: Upstream derivatives.
    - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive

    Returns a tuple of:
    - dx: Gradient with respect to x
    - dw: Gradient with respect to w
    - db: Gradient with respect to b
    """
    dx, dw, db = None, None, None
    N, F, out_height, out_width = dout.shape
    x, w, b, conv_param = cache

    stride, pad = [conv_param['stride'], conv_param['pad']]
    xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
    num_filts, _, f_height, f_width = w.shape

    # ===== #
    # YOUR CODE HERE:
    # Implement the backward pass of a convolutional neural network.
    # Calculate the gradients: dx, dw, and db.

```

```

# ===== #
'''- w: Filter weights of shape (F, C, HH, WW)'''
N, C, H, W = x.shape

db = np.sum(dout, axis=(0,2,3))
dw = np.zeros(w.shape)
dx = np.zeros(x.shape)

mask = np.pad(np.ones(x.shape),((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
for i in range(num_filts):
    for j in range(out_height):
        for k in range(out_width):
            dw[i,:,:,:] += np.sum((dout[:,i,j,k]*xpad[:, :, stride*j:stride*j+f_height, stride*k:stride*k+f_width]), axis=0)

#dx
wrot = np.rot90(w,2,axes=(2,3))
dout_dial = np.zeros((N, F, out_height+(stride-1)*(out_height-1), out_width+(stride-1)*(out_width-1)))
dout_dial[:, :, 0::stride, 0::stride] = dout
_,_,ddp_height,ddp_width = dout_dial.shape
print(dout_dial.shape)
dout_dial_pad = np.pad(dout_dial, ((0,0), (0,0), (H-ddp_height+1,H-ddp_height+1), (W-ddp_width+1,W-ddp_width+1)), mode='constant')

for n in range(N):
    for i in range(num_filts):
        for j in range(H):
            for k in range(W):
                dx[n,:,j,k] += np.sum(wrot[i,:, :, :] * dout_dial_pad[n,i,j:j+f_height,k:k+f_width], axis=0)

# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dw, db

def max_pool_forward_naive(x, pool_param):
    """
    A naive implementation of the forward pass for a max pooling layer.

    Inputs:
    - x: Input data, of shape (N, C, H, W)
    - pool_param: dictionary with the following keys:
      - 'pool_height': The height of each pooling region
      - 'pool_width': The width of each pooling region
      - 'stride': The distance between adjacent pooling regions
    """

```

```

Returns a tuple of:
- out: Output data
- cache: (x, pool_param)
"""
out = None

# ===== #
# YOUR CODE HERE:
# Implement the max pooling forward pass.
# ===== #
ph = pool_param['pool_height']
pw = pool_param['pool_width']
stride = pool_param['stride']
N, C, H, W = x.shape
Hout = int(1 + (H - ph) / stride)
Wout = int(1 + (W - pw) / stride)
out = np.zeros((N,C,Hout,Wout))
for n in range(N):
    for i in range(Hout):
        for j in range(Wout):
            out[n,:,i,j] = np.max(x[n,:,stride*i:stride*i+ph,stride*j:stride*j+pw], axis=(1,2))

# ===== #
# END YOUR CODE HERE
# ===== #
cache = (x, pool_param)
return out, cache

def max_pool_backward_naive(dout, cache):
    """
    A naive implementation of the backward pass for a max pooling layer.

    Inputs:
    - dout: Upstream derivatives
    - cache: A tuple of (x, pool_param) as in the forward pass.

    Returns:
    - dx: Gradient with respect to x
    """
    dx = None
    x, pool_param = cache
    pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'], pool_param['stride']

    # ===== #
    # YOUR CODE HERE:
    # Implement the max pooling backward pass.

```

```

# ===== #
N, C, H, W = x.shape
dx= np.zeros_like(x)
_,_,Hout,Wout = dout.shape
out = np.zeros((N,C,Hout,Wout))
for n in range(N):
    for i in range(Hout):
        for j in range(Wout):
            for k in range(C):
                out[n,k,i,j] = np.max(x[n,k, stride*i:stride*i+pool_height, stride*j:stride*j+pool_width], axis=(0,1))
                mask = (out[n,k,i,j]==x[n,k, stride*i:stride*i+pool_height, stride*j:stride*j+pool_width])
                dx[n,k, stride*i:stride*i+pool_height, stride*j:stride*j+pool_width] = mask*dout[n,k,i,j]
# ===== #
# END YOUR CODE HERE
# ===== #

return dx

def spatial_batchnorm_forward(x, gamma, beta, bn_param):
    """
    Computes the forward pass for spatial batch normalization.

    Inputs:
    - x: Input data of shape (N, C, H, W)
    - gamma: Scale parameter, of shape (C,)
    - beta: Shift parameter, of shape (C,)
    - bn_param: Dictionary with the following keys:
        - mode: 'train' or 'test'; required
        - eps: Constant for numeric stability
        - momentum: Constant for running mean / variance. momentum=0 means that
          old information is discarded completely at every time step, while
          momentum=1 means that new information is never incorporated. The
          default of momentum=0.9 should work well in most situations.
        - running_mean: Array of shape (D,) giving running mean of features
        - running_var Array of shape (D,) giving running variance of features

    Returns a tuple of:
    - out: Output data, of shape (N, C, H, W)
    - cache: Values needed for the backward pass
    """
    out, cache = None, None

    # ===== #
    # YOUR CODE HERE:
    # Implement the spatial batchnorm forward pass.
    #

```

```

# You may find it useful to use the batchnorm forward pass you
# implemented in HW #4.
# ===== #
N,C,H,W = x.shape
x_reshape = np.zeros((N*H*W,C))
for i in range(C):
    x_reshape[:,i] = x[:,i,:,:].reshape(N*H*W)
out_rs,cache =batchnorm_forward(x_reshape, gamma, beta, bn_param)
out = np.zeros_like(x)
for i in range(C):
    out[:,i,:,:) = out_rs[:,i].reshape(N,H,W)
# ===== #
# END YOUR CODE HERE
# ===== #

return out, cache

def spatial_batchnorm_backward(dout, cache):
    """
    Computes the backward pass for spatial batch normalization.

    Inputs:
    - dout: Upstream derivatives, of shape (N, C, H, W)
    - cache: Values from the forward pass

    Returns a tuple of:
    - dx: Gradient with respect to inputs, of shape (N, C, H, W)
    - dgamma: Gradient with respect to scale parameter, of shape (C,)
    - dbeta: Gradient with respect to shift parameter, of shape (C,)
    """
    dx, dgamma, dbeta = None, None, None

    # ===== #
    # YOUR CODE HERE:
    # Implement the spatial batchnorm backward pass.
    #
    # You may find it useful to use the batchnorm forward pass you
    # implemented in HW #4.
    # ===== #
    N,C,H,W = dout.shape
    dout_reshape = np.zeros((N*H*W,C))
    for i in range(C):
        dout_reshape[:,i] = dout[:,i,:,:].reshape(N*H*W)
    dx_rs,dgamma,dbeta = batchnorm_backward(dout_reshape,cache)
    dx = np.zeros_like(dout)

```

```
for i in range(C):
    dx[:,i,:,:] = dx_rs[:,i].reshape(N,H,W)
# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dgamma, dbeta
```